

# Simple features

---

Morten R. Hannemose, [mohan@dtu.dk](mailto:mohan@dtu.dk)

March 8, 2024

02504 Computer vision course lectures,  
DTU Compute, Kgs. Lyngby 2800, Denmark



**This lecture is being  
livestreamed and recorded  
(hopefully)**

# **Two feedback persons**

# Learning objectives

After this lecture you should be able to:

- explain the image correspondence problem
- explain the use of image features (key points, interest points)
- understand and implement image filtering
- implement Harris corner detection
- run Canny edge detection

# Presentation topics

Image correspondence problem

Features and feature descriptors

Simple features

Filtering and convolution

Harris corners

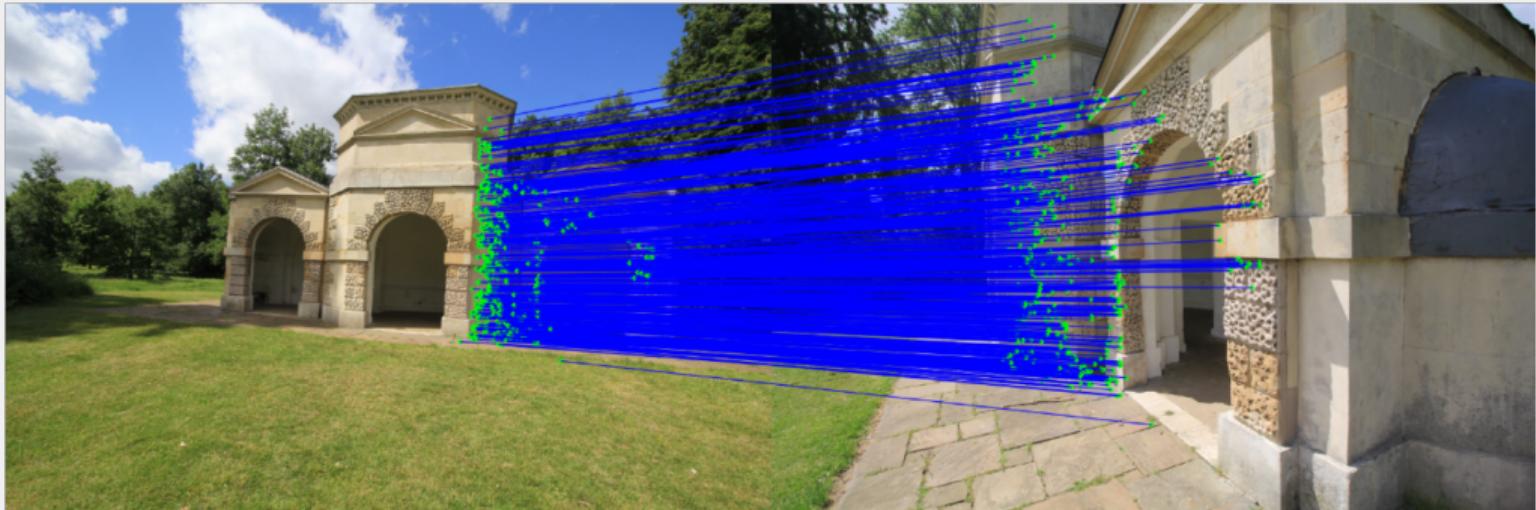
Canny edges

# Image correspondence problem

---

# Image correspondence

The problem of matching two parts of different images:



# Image correspondence

More images of the same scene

- Correspondence exists between parts that are visible in each image
- Not all corresponding points have a unique pattern
- Idea: Only choose points that can be identified uniquely

# What are good points?



# What are good points?

We are able to identify the same corners in multiple images



# Problems

## Movement of camera

- Scale – distance of camera to object
- Rotation – objects and camera is rotated between frames
- Translation – movement of camera from one place to another

Result: Appearance change dependent on distance to camera

# Problems

Movement of camera



# Problems

## Other issues

- Occlusion – objects can be in front of other objects
- Lighting change – darker/lighter, color of light (change in spectrum), direction of light
- Clutter – many objects in a scene

# Solution

## Invariance to imaging problems

- Focus on points and a small area around each point
- Two aspects
  - identify key points (focus of today)
  - characterize pattern around point (week 8)

# Features and feature descriptors

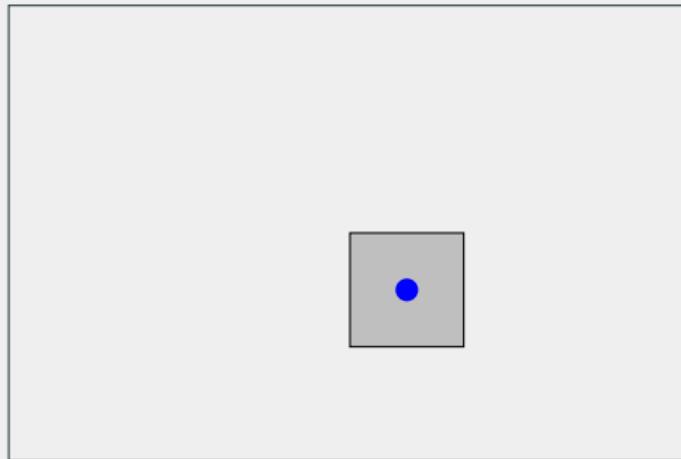
First some terminology:

Key points, interest points, and feature points are the same – namely points in an image with a coordinate position (also between pixels).

Descriptors, key point descriptors, interest point descriptors, and feature descriptors are the same and used for characterizing the pattern around a point. Usually a vector that can be matched between images.

# The image patch feature descriptor

One easy way of describing a feature is to use the local pixel values around the feature.



## The image patch matching

Stretch the image patch into a vector  $f$ , and that is your descriptor.

A simple comparison operator  $d(f_1, f_2)$  just uses the scalar product as the distance between features:

$$d(f_1, f_2) = f_1^T f_2$$

Good with only small transformations, low noise, and unchanged environment.

## Cross correlation

A better comparison  $X(\mathbf{f}_1, \mathbf{f}_2)$  uses the **cross correlation** as the distance:

$$X(\mathbf{f}_1, \mathbf{f}_2) = \frac{\text{cov}(\mathbf{f}_1, \mathbf{f}_2)}{\sqrt{\text{var}(\mathbf{f}_1) + \text{var}(\mathbf{f}_2)}}, \text{ where}$$

$$\text{cov}(\mathbf{p}, \mathbf{q}) = \frac{1}{N-1} \sum_i^N (\mathbf{p}_i - \bar{\mathbf{p}})^T (\mathbf{q}_i - \bar{\mathbf{q}}), \text{ and}$$

$$\text{var}(\mathbf{p}) = \frac{1}{N-1} \sum_i^N \|\mathbf{p}_i - \bar{\mathbf{p}}\|^2.$$

## Cross correlation

- The cross correlation picks up local variation, and is therefore more robust against environment changes than squared distance
- Only useful when images are extremely similar.

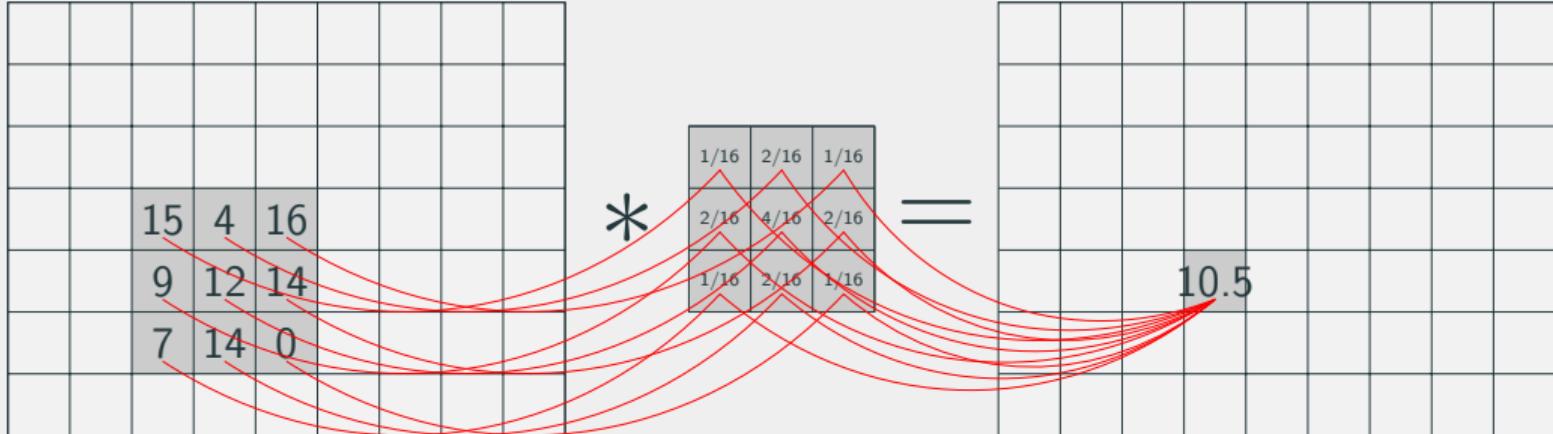
More robust (but more complicated) feature descriptors will follow in two weeks.

# Simple features

---

# Filtering and convolution

$$I(x; t) = \int_{\xi \in \mathbb{R}^2} I(x - \xi) g(\xi; t) d\xi \approx \sum_{i=-k}^k \sum_{j=-k}^k I(x - i, y - j) g(i, j)$$



## 2D convolution

We use filtering and convolution for the same operation and use the symbol  $*$  for convolution. Convolution is commutative

$$I_g = g * I = I * g$$

2D Gaussian  $g : \mathbb{R}^2 \times \mathbb{R}_+ \rightarrow \mathbb{R}$

$$g(x, y; \sigma^2) = \frac{1}{2\pi\sigma^2} \exp\left(\frac{-(x^2 + y^2)}{2\sigma^2}\right). \quad (1)$$

## Convolution – separability

Useful property: The (isotropic) Gaussian is separable

$$g(x, y; \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-y^2}{2\sigma^2}\right)$$

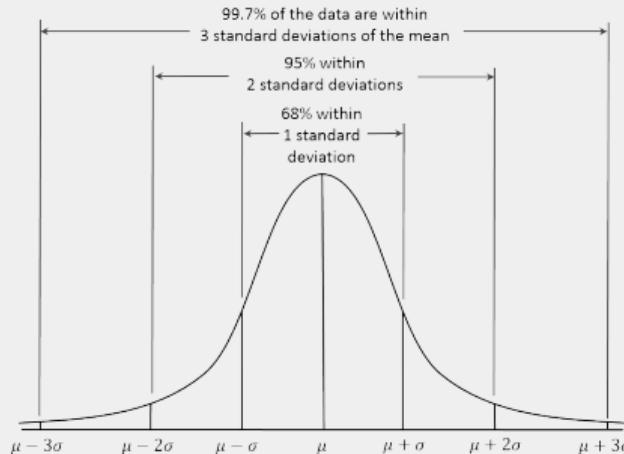
This separability means that

$$I_g = (\mathbf{g} * \mathbf{g}^T) * I = \mathbf{g} * (\mathbf{g}^T * I)$$

where  $\mathbf{g}$  is a vector of the Gaussian.

# Size of Gaussian filter

- Ideally: Infinitely big – but this is impractical
- Empiric rule:  $3\sigma$ ,  $4\sigma$  or  $5\sigma$  rule – example:
  - if  $\sigma = 2$ , filter size of Gaussian is  $2 \cdot 5 \cdot 2 + 1 = 21$
  - if  $\sigma = 20$ , filter size of Gaussian is  $2 \cdot 5 \cdot 20 + 1 = 201$



## **Demo – Gaussian filtering**

`cv2.filter2D`

`cv2.sepFilter2D`

# Derivatives of an image

How can we find the derivative of an image in e.g. the  $x$ -direction?  
What does it mean to take a derivative?

# Derivatives of an image

How can we find the derivative of an image in e.g. the  $x$ -direction?

What does it mean to take a derivative?

An image is only discrete values.

If we can make the image **continuous** all our math works again

## Derivatives of an image using Gaussians

Let  $\mathbf{g}$  be column vector of a Gaussian. Consider the blurred image

$$I_b = \mathbf{g} * \mathbf{g}^T I.$$

## Derivatives of an image using Gaussians

Let  $\mathbf{g}$  be column vector of a Gaussian. Consider the blurred image

$$I_b = \mathbf{g} * \mathbf{g}^\top I.$$

Then the derivative of  $I_b$  in the  $x$ -direction is

$$\begin{aligned}\frac{\partial}{\partial x} I_b &= \frac{\partial}{\partial x} (\mathbf{g} * \mathbf{g}^\top * I) \\ &= \mathbf{g} * \left( \frac{\partial}{\partial x} \mathbf{g}^\top \right) * I \\ &= \mathbf{g} * \mathbf{g}_d^\top * I,\end{aligned}$$

# Derivative of the Gaussian

Recall the one-dimensional Gaussian

$$g(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-x^2}{2\sigma^2}\right).$$

We can then compute the derivative

$$g_d(x) = \frac{d}{dx}g(x) = \frac{-x}{\sigma^2}g(x).$$

This is straightforward to implement on a computer.

# **Short break**

## Harris corners

Corners (and blobs) are great features because they are easy to describe and detect.

Let's see how much the intensity changes, for a small shift  $\Delta_x, \Delta_y$

$$\Delta I(x, y, \Delta_x, \Delta_y) = I(x, y) - I(x + \Delta_x, y + \Delta_y)$$

Harris corners are detected as points with *locally maximum* change from a small shift.

A corner is then a local area where  $\Delta I(x, y, \Delta_x, \Delta_y)^2$  is large no matter how we choose  $\Delta_x, \Delta_y$ .

## Harris corner measure

Because the “cornerness” should be true for a local area and not just a point, we apply another Gaussian blur to the above measure.

The measure to check for corners is then

$$\begin{aligned} c(x, y, \Delta_x, \Delta_y) &= g * \Delta I(x, y, \Delta_x, \Delta_y)^2 \\ &= g * (I(x, y) - I(x + \Delta_x, y + \Delta_y))^2, \end{aligned}$$

where  $g * \dots$  is the convolution with the Gaussian.

## Replace with Taylor approximation

$$I(x + \Delta_x, y + \Delta_y) \approx I(x, y) - \frac{\partial I}{\partial x}(x, y) \Delta_x - \frac{\partial I}{\partial y}(x, y) \Delta_y$$

## Replace with Taylor approximation

$$I(x + \Delta_x, y + \Delta_y) \approx I(x, y) - \underbrace{\frac{\partial I}{\partial x}(x, y) \Delta_x}_{I_x} - \underbrace{\frac{\partial I}{\partial y}(x, y) \Delta_y}_{I_y}$$

## Replace with Taylor approximation

$$\begin{aligned} I(x + \Delta_x, y + \Delta_y) &\approx I(x, y) - \underbrace{\frac{\partial I}{\partial x}(x, y) \Delta_x}_{I_x} - \underbrace{\frac{\partial I}{\partial y}(x, y) \Delta_y}_{I_y} \\ &= I(x, y) - [I_x \quad I_y] \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix} \end{aligned}$$

$I_x$  and  $I_y$  are also values at  $(x, y)$ , but we omit it for readability.

## Harris corner measure derivation

The corner checking measure is then

$$\begin{aligned} c(x, y, \Delta_x, \Delta_y) &= g * \left( I(x, y) - I(x + \Delta_x, y + \Delta_y) \right)^2 \\ &\approx g * \left( \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix} \right)^2 \end{aligned}$$

## Harris corner measure derivation

The corner checking measure is then

$$\begin{aligned} c(x, y, \Delta_x, \Delta_y) &= g * \left( I(x, y) - I(x + \Delta_x, y + \Delta_y) \right)^2 \\ &\approx g * \left( \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix} \right)^2 \\ &= g * \left( \begin{bmatrix} \Delta_x & \Delta_y \end{bmatrix} \begin{bmatrix} I_x \\ I_y \end{bmatrix} \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix} \right) \\ &= g * \left( \begin{bmatrix} \Delta_x & \Delta_y \end{bmatrix} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix} \right) \end{aligned}$$

## Harris corner measure derivation

Finally, we take the convolution with the Gaussian inside

$$\begin{aligned} c(x, y, \Delta_x, \Delta_y) &\approx g * \left( [\Delta_x \quad \Delta_y] \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix} \right) \\ &= [\Delta_x \quad \Delta_y] \underbrace{\begin{bmatrix} g * (I_x^2) & g * (I_x I_y) \\ g * (I_x I_y) & g * (I_y^2) \end{bmatrix}}_{C(x,y)} \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix}. \end{aligned}$$

We end up with  $C(x, y)$ .

## The structure tensor

Corners have a large  $c(x, y, \Delta_x, \Delta_y)$  regardless of  $[\Delta_x \ \Delta_y]$ .

$$c(x, y, \Delta_x, \Delta_y) = [\Delta_x \ \Delta_y] \mathbf{C}(x, y) \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix}$$

How can we check if this the case from the  $\mathbf{C}(x, y)$  matrix?

## The structure tensor

Corners have a large  $c(x, y, \Delta_x, \Delta_y)$  regardless of  $[\Delta_x \ \Delta_y]$ .

$$c(x, y, \Delta_x, \Delta_y) = [\Delta_x \ \Delta_y] \mathbf{C}(x, y) \begin{bmatrix} \Delta_x \\ \Delta_y \end{bmatrix}$$

How can we check if this the case from the  $\mathbf{C}(x, y)$  matrix?

When both eigenvalues of  $\mathbf{C}(x, y)$  are large.

## The Harris corner metric continued

Let  $\lambda_1$  and  $\lambda_2$  be the eigenvalues of  $\mathbf{C}(x, y) = \begin{bmatrix} a & c \\ c & b \end{bmatrix}$ .

We then have the following

$$\det(\mathbf{C}(x, y)) = \lambda_1 \lambda_2 = ab - c^2,$$

$$\text{trace}(\mathbf{C}(x, y)) = \lambda_1 + \lambda_2 = a + b, \text{ let us then define}$$

$$r(x, y) = \det(\mathbf{C}(x, y)) - k \text{ trace}(\mathbf{C}(x, y))^2$$

## The Harris corner metric continued

Let  $\lambda_1$  and  $\lambda_2$  be the eigenvalues of  $\mathbf{C}(x, y) = \begin{bmatrix} a & c \\ c & b \end{bmatrix}$ .

We then have the following

$$\det(\mathbf{C}(x, y)) = \lambda_1\lambda_2 = ab - c^2,$$

$$\text{trace}(\mathbf{C}(x, y)) = \lambda_1 + \lambda_2 = a + b, \text{ let us then define}$$

$$\begin{aligned} r(x, y) &= \det(\mathbf{C}(x, y)) - k \text{ trace}(\mathbf{C}(x, y))^2 \\ &= \lambda_1\lambda_2 - k(\lambda_1 + \lambda_2)^2 \\ &= ab - c^2 - k(a + b)^2 \end{aligned}$$

# The Harris corner metric

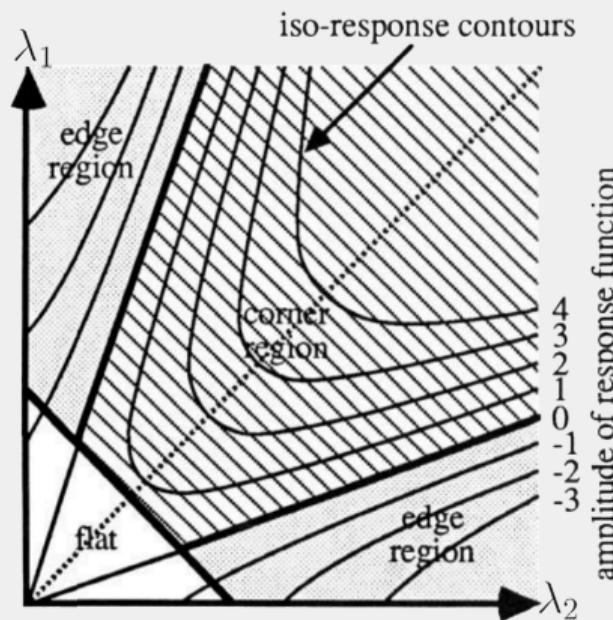
Now we have the Harris corner metric

$$r(x, y) = \lambda_1\lambda_2 - k(\lambda_1 + \lambda_2)^2.$$

$k$  is a free parameter, typically  $k = 0.06$ .

Notice that  $r(x, y)$  is:

- negative for one eigenvalue much greater than the other
- large and positive for large eigenvalues
- small and positive for small eigenvalues



## The Harris corner detector

We can now detect corners by finding points where  $r(x, y)$  is greater than some threshold  $\tau$ .

Typically, you can choose  $\tau$  to be

$$0.1 \cdot \max(r(x, y)) < \tau < 0.8 \cdot \max(r(x, y)).$$

## Non-maximum suppression

Find the local maximum of one pixel compared to neighbours

$$(I(x, y) - I(x', y')) > 0 \quad \forall x' \in n(x, y)$$

where  $n(x, y)$  is a neighbourhood around the point  $(x, y)$

Suggested procedure

- Initialize an array of size of image with  $I(x, y) > \tau$
- Compare entire image with one neighbour (e.g to the right)
- Set pixels to 0 where it is not larger than neighbour
  - Repeat for all neighbours
- Find coordinates of all pixels that are 1

## Canny edges

Often, we might also consider detecting lines.

To detect (non-straight) lines, we can use the Canny edge detector.

## Canny edges

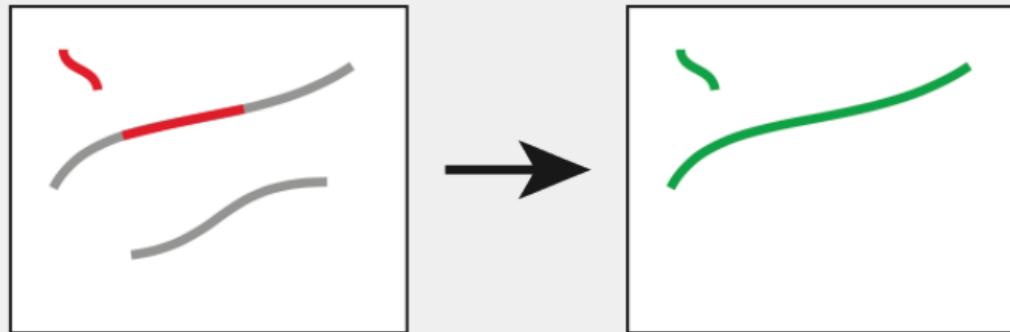
The metric in the canny edge detector is simply the gradient magnitude

$$m(x, y) = \sqrt{I_x^2(x, y) + I_y^2(x, y)}.$$

And the edges are detected by thresholding the magnitude  $m(x, y)$ , however, in two stages: **seed and grow**.

# Canny edges

- Seed: label edges by a threshold where  $m(x, y) > \tau_1$
- Grow: with a second threshold where  $m(x, y) > \tau_2$ , label iff the new points are next to previously labelled edges.
- The threshold values are chosen such that  $\tau_1 > \tau_2$



# Exercise

During the exercise, you will

- **implement** the Harris corner detector
- **try** the Canny edge detector

**Suggestion:** After finishing your implementation – try on your own images and perhaps draw some features and see how it works

# Learning objectives

After this lecture you should be able to:

- explain the image correspondence problem
- explain the use of image features (key points, interest points)
- understand and implement image filtering
- implement Harris corner detection
- run Canny edge detection

**Exercise time!**